Scheduling Assistant: Visual Output Options

For the client demo build | Prepared by Nisha | 2 July 2026

This document covers the three standard ways manufacturing companies visualise production schedules, the idea behind each, and how each one fits into our demo build. The guiding principle: factory teams already have mental models for schedules. We digitise a format they recognise instead of inventing a new one. The AI layer, not the chart, is what makes it modern.

Option 1: The machine loading board (Gantt view)

The idea. This is the default view in production planning worldwide. Machines or lines run down the side, time runs across the top, and each order is a block. Plant teams have seen this on paper and whiteboards for decades, so it needs zero training. Yes, it is old school. For this audience, old school means instantly familiar, and familiarity drives adoption.

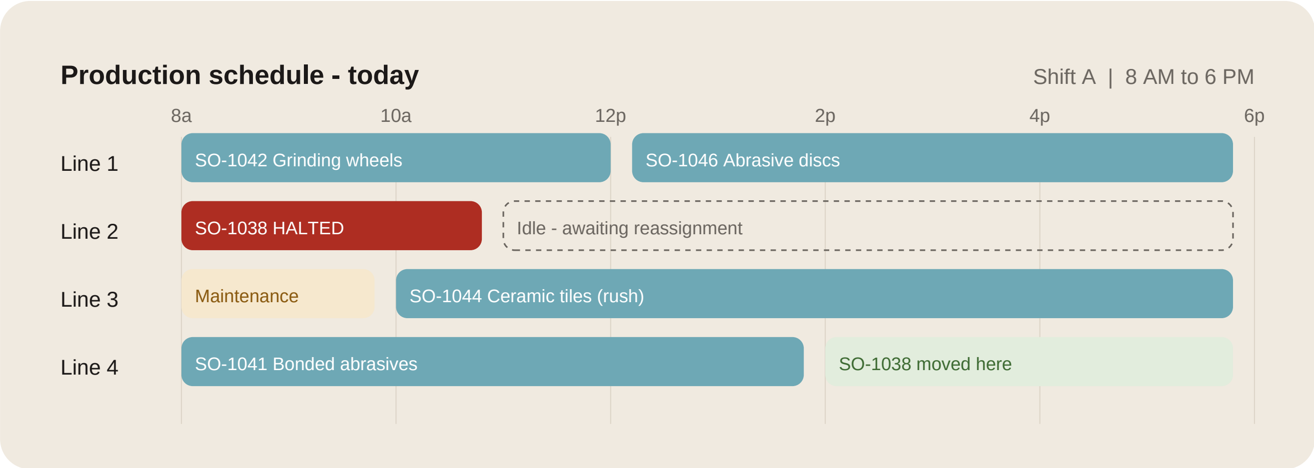

AI note: Line 2 spindle fault at 10:40. SO-1038 moved to Line 4 after SO-1041. New delivery date: 9 Jul (was 8 Jul). Sales team notified.

How Divya can apply it

- This is the main screen of the scheduling assistant. One row per production line, one coloured block per order.
- Each block needs five pieces of dummy data: order number, product name, assigned line, start time, end time.
- Add one status per block: running (teal), halted (red), maintenance (amber), rescheduled (green). Colour is the status.
- Demo moment: trigger a machine breakdown. The halted block turns red, the system moves the order to a free line, and the AI note below the chart explains what changed and the new delivery date. This covers all three of Anusha's scenarios: delay, breakdown, urgent order.
- Keep it read-only for the demo. No drag and drop needed. The AI does the moving; the human just sees the result.

Option 2: The daily run sheet (dispatch list)

The idea. This is the single Excel-looking sheet that gets printed and pasted next to the machine every morning. It answers one question: what do we run today, and in what order. Every supervisor reads this format without being taught. In our demo it becomes the output of the assistant: someone asks a question, and the answer arrives as an updated run sheet.

Seq	Order	Product	Line	Due	Status
1	SO-1044	Ceramic tiles	3	3 Jul	RUSH
2	SO-1042	Grinding wheels	1	5 Jul	RUNNING
3	SO-1038	Bonded abrasives	4	9 Jul	RESCHEDULED
4	SO-1046	Abrasive discs	1	11 Jul	ON TRACK

How Divya can apply it

- Build this as a simple table view that sits one click away from the Gantt board. Same data, different shape.
- Columns: sequence number, order number, product, line, due date, status. Nothing more. Six columns maximum.
- Add a print button. The printable version is the adoption story: it looks exactly like the paper sheet they use today, except the AI keeps it current.
- When the breakdown scenario runs on the Gantt board, this sheet should re-sequence automatically. Showing both views update from one event is the strongest moment in the demo.

Option 3: The order milestone tracker

The idea. This one is not for the shop floor. It is for the salesperson standing in front of a customer asking where the order is. Instead of calling the project manager, they ask the chatbot, and the answer comes back as a milestone strip: which stages are done, where the order is stuck, and what to do about it. Think of a courier tracking page, applied to a production order. This is the visual that connects the knowledge center and the scheduling assistant into one story.

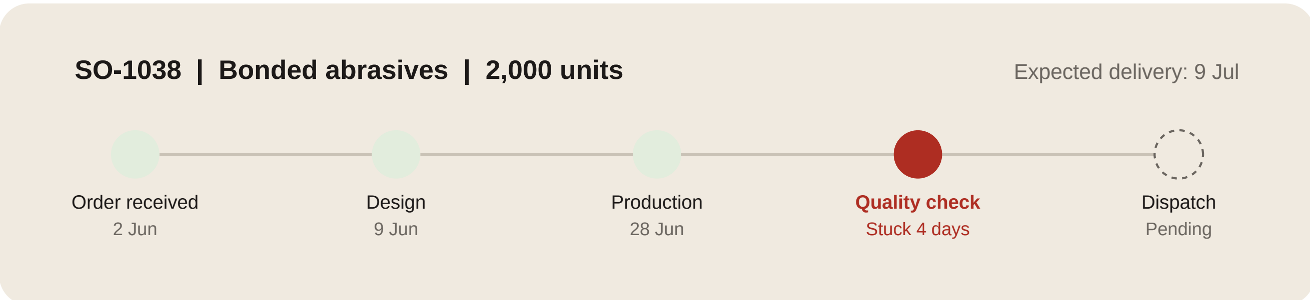

AI note: Held at quality check pending supervisor sign-off since 28 Jun. Suggested action: escalate to QC lead. Delivery date at risk if not cleared by 5 Jul.

How Divya can apply it

- This lives inside the chatbot, not as a separate screen. When a user asks about a specific order number, the reply includes this strip.
- Five fixed stages work for the demo: order received, design, production, quality check, dispatch. Each stage needs a date or a pending flag.
- Add one stuck rule: if an order sits in a stage beyond a set number of days, mark it red and show a suggested action. That single rule demonstrates the flagging behaviour Anusha described.
- Because this appears in the chatbot, it quietly proves the point we agreed on the call: both systems draw from the same data source, even though we present them as two products.

Recommendation: layer all three

Do not pick one. Each visual serves a different person in the room, and together they cover the full demo narrative:

- **Machine loading board** is the planner's main screen. This is where the breakdown scenario plays out live.
- **Daily run sheet** is the shop floor output. Printable, familiar, zero learning curve. This answers the adoption concern directly.
- **Milestone tracker** is the chatbot's answer format for order queries. This is the sales and customer-facing story.

One dummy dataset feeds all three views. Roughly six to eight orders, four production lines, one shift, one breakdown event. Build the dataset first, then the three views on top of it. If time runs short before the demo, the priority order is: Gantt board first, milestone tracker second, run sheet third.

The dummy dataset checklist

Every order in the dataset needs: order number, product name, quantity, assigned line, start time, end time, current stage, stage entry date, due date, and status. Ten fields, six to eight orders. That is the entire data requirement for the scheduling assistant demo.

Open decision from the 2 July call: the team has not locked the visualisation format. This document is the input for that decision. Target: lock it early in the week of 7 July so the build is not blocked.